MODULE *PortAllocator*

The port allocator updates an endpoint's ports field from its spec. For host-mode ports, it just copies them over. For *ingress*-mode ports, it checks that the requested port number is free (if given), or allocates an unused one from the dynamic range (if not).

Places where I deviated from the Go code to make the checks pass are marked with *XXX*. In summary:

Reusing a previous assignment when we need it elsewhere

Example:

1. [*name* = "foo", dynamic] (initial configuration)
2. [*name* = "foo", dynamic], [*name* = "bar", published = 30000] (updated)

Here we will reject the update because we try to reuse port 30000 for "foo", even though we now need it for "bar".

Duplicate similar ports

Example:

1. [*name* = "foo", dynamic]
2. [*name* = "foo", dynamic], [*name* = "foo", dynamic]

We will try to use port number 30000 for both ports and reject the allocation.

Shadowing an existing host port

The allocator ignores host ports completely. It may allocate a dynamic *ingress* port using a network address that is already in use on some host. In this case, the original service becomes unreachable.

Accepting a host port that is hidden

The allocator may accept a host-port allocation that it knows will be unreachable on any host because that address is already in use as an *ingress* port.

EXTENDS *Sequences*, *Integers*, *TLC*, *FiniteSets* Some libraries we use

$Range(S) \triangleq \{S[i] : i \in \text{DOMAIN } S\}$ Generic helper function

The set of protocols we support (*e.g.* {"*tcp*", "*udp*", "*sctp*"}). We assume that every protocol has the same set of static and dynamic ports.

CONSTANT *Protocol*

A host-mode port is available only via the host that ends up running the workload. Two different services can use the same port, as long as they run on different hosts.

$host \triangleq$ "host"

Ingress-mode ports are available via any node in the cluster. Connections will be forwarded to a host running the service.

$ingress \triangleq$ "ingress"

$Mode \triangleq \{ingress, host\}$

Port numbers that can be assigned by the allocator when the user requests a dynamic port number. The code currently uses 30000 . . 32767.

CONSTANT $DynamicPortNumber$

Non-dynamic port numbers.
CONSTANT $StaticPortNumber$

All port numbers. *SwarmKit* uses 1 .. 65535. We mainly define it this may to make it easy to make dynamic and static numbers symmetry sets in the model checker.
$PortNumber \triangleq DynamicPortNumber \cup StaticPortNumber$

A special value for the *published_port* field to indicate that the allocator should select the port. *SwarmKit* uses 0 for this.
$dynamic \triangleq \text{CHOOSE } x : x \notin PortNumber$

The type of endpoint *IDs* (*e.g.* STRING).
CONSTANT $EndpointID$

The type of port names (*e.g.* STRING).
CONSTANT $Name$

The maximum number of ports in a specification.
CONSTANT $MaxPorts$

The requirements for a port, provided by the user. The user can specify a port number directly (including any number in *DynamicPortNumber*), or can specify *dynamic* to have the system allocate it.

The real structure also includes *target_port* , which is the port inside the container. For the model, we can consider this as part of *name* (it just makes the port more unique).

$$
\begin{array}{l}
PortSpec \triangleq [\\
\quad mode : Mode, \\
\quad name : Name, \\
\quad protocol : Protocol, \\
\quad published\_port : StaticPortNumber \cup DynamicPortNumber \cup \{dynamic\} \\
]
\end{array}
$$

A configured port, after the allocator has done its job. Note: The *SwarmKit* code uses a single Go type for this and *PortSpec*.

$$
\begin{array}{l}
PortConfig \triangleq \\
\quad \text{Either an } ingress \text{ port:} \\
\quad [\\
\quad\quad mode : \{ingress\}, \\
\quad\quad name : Name, \\
\quad\quad protocol : Protocol, \\
\quad\quad published\_port : PortNumber \quad \text{The port is now allocated} \\
\quad] \\
\quad \cup \quad \text{or a } host \text{ port:} \\
\quad [\\
\quad\quad mode : \{host\}, \\
\quad\quad name : Name,
\end{array}
$$

```
    protocol       : Protocol,
    published_port : PortNumber ∪ {dynamic}   Can still be unassigned
  ]

  Two PortConfig/PortSpec values are "mostly equal" if they differ only in their published_port.
PortsMostlyEqual(a, b) ≜
  LET ignorePP(x) ≜ [f ∈ DOMAIN x \ {"published_port"} ↦ x[f]]
  IN  ignorePP(a) = ignorePP(b)

  A network address is a protocol and port-number pair.
Address ≜ Protocol × PortNumber

  The network address of a SwarmKit port.
Addr(port) ≜
  ⟨port.protocol, port.published_port⟩

  A finite sequence of maximum length max . Useful for model checking.
FiniteSeq(S, max) ≜
  UNION {[1 .. n → S] : n ∈ 0 .. max}

  An endpoint specification is just a list of port specifications.
EndpointSpec ≜ FiniteSeq(PortSpec, MaxPorts)

  An endpoint object records the currently allocated ports and the original specification that led
  to this assignment.
Endpoint ≜ [
  spec  : EndpointSpec,
  ports : Seq(PortConfig)
]

  nullEndpoint is used to represent an endpoint that does not yet exist.
nullEndpoint ≜ [
  spec  ↦ ⟨⟩,
  ports ↦ ⟨⟩
]

  The allocator returns a proposal for updating the state, rather than doing it immediately. Note
  that the port allocator's proposal is only valid until the next time an allocation is requested.

  The allocate field is not really needed here, as we can generate it easily (it's just the set of
  ingress addresses in ports ). deallocate wouldn't be needed if the commit operation took the
  old configuration as an argument, as it's just the ingress addresses in that. However, the Go
  code includes these fields, so we do too.
Proposal ≜ [
  deallocate : SUBSET Address,   Addresses to remove from allocated
  allocate   : SUBSET Address,   Addresses to add to allocated (after deallocation)
  ports      : Seq(PortConfig)   The new value for endpoint.ports
]
```

The allocator

The set of allocated *ingress* addresses.

VARIABLE *allocated*

The smallest item in a non-empty set.

$MinElement(S) \triangleq$ CHOOSE $x \in S : \forall\, y \in S : x \leq y$

Return an updated version of *spec* in which dynamic *ingress* ports have been updated to copy the existing configution, where possible.

$RecoverExistingPorts(endpoint,\ spec) \triangleq$

LET The current configuration

$oldPorts \triangleq endpoint.ports$

Indexes in the list of ports for which we need to choose a port number:

$dynamics \triangleq \{i \in$ DOMAIN $spec :$
$\wedge\ spec[i].mode = ingress$
$\wedge\ spec[i].published\_port = dynamic\}$

Ingress ports for which the user specified the port:

$forcedPorts \triangleq \{p \in Range(spec) :$
$\wedge\ p.mode = ingress$
$\wedge\ p.published\_port \neq dynamic\}$

The (*ingress*) addresses the user specified manually:

$forcedAddrs \triangleq \{Addr(p) : p \in forcedPorts\}$

The recovered port number for port i :

$Recover(i) \triangleq$

LET The port specification from the user:

$s \triangleq spec[i]$

The currently configured $port(s)$ that are like s :

$olds \triangleq \{j \in$ DOMAIN $oldPorts :$
$\wedge\ \exists\, k \in$ DOMAIN $endpoint.spec :$
$\wedge\ endpoint.spec[k] = s$ Spec hasn't changed
$\wedge\ PortsMostlyEqual(oldPorts[j],\ s)$
We're not forced to use this for something else:
XXX: does *SwarmKit* do this?
$\wedge\ Addr(oldPorts[j]) \notin forcedAddrs$
$\}$

Whether s is similar to a previous dynamic port in the spec list:

XXX: Looks like the *SwarmKit* code doesn't do this check.

$duplicate \triangleq \exists j \in 1\,..\,(i-1) :$ An earlier similar port in the list
$\wedge\ j \in dynamics$ needed a dynamic assignment too
$\wedge\ PortsMostlyEqual(spec[i],\ spec[j])$

IN

IF $\vee\ olds = \{\}$ If we haven't already got anything like s
$\vee\ duplicate$ Or we already used it

```
        THEN dynamic  Then don't update s − it's still dynamic
        ELSE
          Use the first of the candidates for published_port :
          oldPorts[MinElement(olds)].published_port
      The updates to apply:
    recovered ≜ [i ∈ dynamics ↦ [spec[i] EXCEPT !.published_port = Recover(i)]]
  IN
  recovered @@ spec  Combine updates with other entries
```

Allocate(*endpoint*, *spec*) returns a set of possible proposals to update *endpoint* to *spec* .

The real system will only return a single proposal. For cases where this function returns {} a real implementation must reject the request. For cases where it returns a non-empty set, a real implementation must return one of the elements as its proposal.

```
Allocate(endpoint, specFromUser) ≜
  LET  Step 1 : Recover dynamic ports from old configuration
       spec ≜ RecoverExistingPorts(endpoint, specFromUser)
       Step 2 : Reject bad user requests due to static assignments
       All the ingress ports in the existing configuration:
       oldIngressPorts ≜ {p ∈ Range(endpoint.ports) : p.mode = ingress}
       Addresses currently in use by endpoint :
       deallocate ≜ {Addr(p) : p ∈ oldIngressPorts}
       Addresses used by other endpoints:
       addrsForOthers ≜ allocated \ deallocate
       Did the user request a port that another endpoint is using?
       alreadyInUse ≜ ∃ p ∈ Range(spec) :
                          ∧ p.mode = ingress
                          ∧ p.published_port ≠ dynamic
                          ∧ Addr(p) ∈ addrsForOthers
       Did the user specify the same static (ingress) address twice?
       haveForcedDuplicates ≜
         ∃ i, j ∈ DOMAIN spec :
           ∧ i ≠ j
           ∧ LET si ≜ spec[i]
                 sj ≜ spec[j]
             IN
               ∧ si.mode = ingress ∧ si.published_port ≠ dynamic
               ∧ sj.mode = ingress ∧ sj.published_port ≠ dynamic
               ∧ Addr(si) = Addr(sj)
  IN
  IF alreadyInUse ∨ haveForcedDuplicates THEN {}  Reject
   ELSE
   Step 3 : Assign dynamic ports
```

There are various ways of assigning the ports. *e.g.* picking the lowest free port, starting the search from the last allocated number, checking already-free ports first and then using ports from *deallocate* only as a last resort. We'll avoid over-specifying by allowing any behaviour here.

```
LET
    Ingress ports that still need to be assigned:
  portsNeeded ≜ {i ∈ DOMAIN spec :
                    ∧ spec[i].mode = ingress
                    ∧ spec[i].published_port = dynamic}
  Possible ways of allocating them. Each element of this set is a mapping from a port index
  to a port number in the dynamic range.
  allocs ≜
    {alloc ∈ [portsNeeded → DynamicPortNumber] :
      Check that alloc is reasonable:
     LET NA(i) ≜ The proposed network address of i
       IF i ∈ DOMAIN alloc THEN ⟨spec[i].protocol, alloc[i]⟩
        ELSE Addr(spec[i])
     IN
     ∀ i ∈ DOMAIN alloc : For each dynamic ingress port:
        No other endpoint is using this address:
       ∧ NA(i) ∉ addrsForOthers
        We're not already trying to allocate this address:
       ∧ ∀ j ∈ DOMAIN spec \ {i} :
            ∨ spec[j].mode = host
            ∨ NA(i) ≠ NA(j)
    }
   Create a proposal object from an allocation mapping:
  Result(alloc) ≜
   LET
     ports ≜ [i ∈ DOMAIN alloc ↦
                [spec[i] EXCEPT !.published_port = alloc[i]]
              ] @@ spec
     ingressPorts ≜ {p ∈ Range(ports) : p.mode = ingress}
   IN
   [
     deallocate ↦ deallocate,
     allocate ↦ {Addr(p) : p ∈ ingressPorts},
     ports ↦ ports
   ]
IN
{Result(x) : x ∈ allocs}

The result of applying prop to the current allocations.
Apply(prop) ≜
  (allocated \ prop.deallocate) ∪ prop.allocate
```

The test system (allocator + user)

The set of active *endpoints* (the allocator doesn't look at this)

VARIABLES *endpoints*

$vars \triangleq \langle allocated,\ endpoints \rangle$

The user creates a new endpoint

$NewEndpoint \triangleq$

$\exists\, s \in EndpointSpec :$ — s is the new spec

$\exists\, id \in EndpointID \setminus \text{DOMAIN } endpoints :$ — id is an unused endpoint ID

$\exists\, prop \in Allocate(nullEndpoint,\ s) :$ — $prop$ is a proposal from the allocator

LET $e \triangleq [spec \mapsto s,\ ports \mapsto prop.ports]$ — e is the new endpoint

IN — Update the store:

$\land\ endpoints' = id :> e \mathbin{@@} endpoints$ — Add e to $endpoints$

$\land\ allocated' = Apply(prop)$ — Tell the allocator to commit

The user updates an existing endpoint

$UpdateEndpoint \triangleq$

$\exists\, s \in EndpointSpec :$ — s is the new spec

$\exists\, id \in \text{DOMAIN } endpoints :$ — id is an existing endpoint

$\exists\, prop \in Allocate(endpoints[id],\ s) :$ — $prop$ is a proposal from the allocator

LET $e \triangleq [spec \mapsto s,\ ports \mapsto prop.ports]$ — e is the new endpoint

IN

$\land\ endpoints' = [endpoints \text{ EXCEPT } ![id] = e]$

$\land\ allocated' = Apply(prop)$

Remove an existing endpoint

$RemoveEndpoint \triangleq$

$\exists\, id \in \text{DOMAIN } endpoints :$ — id is an existing endpoint

LET $props \triangleq Allocate(endpoints[id],\ \langle\rangle)$ — Ask the allocator to remove all ports

IN

$\land\ Assert(props \neq \{\},$ "Rejected remove operation!")

$\land \exists\, prop \in props :$

Commit the removal proposal

$\land\ endpoints' = [i \in \text{DOMAIN } endpoints \setminus \{id\} \mapsto endpoints[i]]$

$\land\ allocated' = Apply(prop)$

The initial state of the system, with no *endpoints* or allocations. When restarting *SwarmKit*, saved *endpoints* can be loaded and allocated as if they were being added as new services using the *Restore* operation. Note: *SwarmKit* does not check whether the saved allocations are consistent at restore time.

$Init \triangleq$

$\land\ endpoints = \langle\rangle$

$\land\ allocated = \{\}$

The possible ways of using the allocator.

$Next \triangleq$

$\lor\ NewEndpoint$

$\lor\ UpdateEndpoint$

$\lor$ *RemoveEndpoint*

Spec $\triangleq$
$\quad$ *Init* $\land \Box[Next]_{vars}$

Properties to check

Check that the variables have the expected types.

TypeOK $\triangleq$
$\quad \land$ *allocated* $\subseteq$ *Address* — A set of addresses
$\quad \land$ DOMAIN *endpoints* $\subseteq$ *EndpointID* — A partial map from endpoint *IDs*
$\quad \land$ *endpoints* $\in$ [DOMAIN *endpoints* $\to$ *Endpoint*] — to *Endpoints*.

Check that the state of the system is consistent: all addresses marked as allocated are needed by some endpoint, all *endpoints* have a configuration that matches their requirements, and no two *endpoints* have been allocated the same address.

AllocationsOK $\triangleq$
$\quad$ Every port the allocator thinks is allocated is actually used by some endpoint
$\quad \land \forall$ *addr* $\in$ *allocated* :
$\quad\quad \exists\, e \in Range(endpoints)$:
$\quad\quad \exists\, p \in Range(e.ports)$:
$\quad\quad Addr(p) = addr$
$\quad$ Every endpoint's configuration is correct
$\quad \land \forall$ *eid* $\in$ DOMAIN *endpoints* :
$\quad\quad$ LET *e* $\triangleq$ *endpoints*[*eid*]
$\quad\quad$ IN
$\quad\quad \land$ *Len*(*e.spec*) = *Len*(*e.ports*) — We have the right number of ports configured
$\quad\quad \land \forall\, i \in$ DOMAIN *e.spec* : — For each port ...
$\quad\quad\quad$ LET *spec* $\triangleq$ *e.spec*[*i*]
$\quad\quad\quad\quad$ *port* $\triangleq$ *e.ports*[*i*]
$\quad\quad\quad$ IN
$\quad\quad\quad$ The actual port is the same as its specification, ignoring dynamic *ingress* port numbers.
$\quad\quad\quad \land$ IF *spec.mode* = *ingress* $\land$ *spec.published_port* = *dynamic*
$\quad\quad\quad\quad$ THEN *PortsMostlyEqual*(*spec*, *port*)
$\quad\quad\quad\quad$ ELSE *spec* = *port*
$\quad\quad\quad$ The port's address is in the *allocated* set.
$\quad\quad\quad \land$ *port.mode* = *ingress* $\Rightarrow$ *Addr*(*port*) $\in$ *allocated*

There are no other users of this port. We only check *spec.mode* = *ingress* because we don't check collisions between host ports here and we'll find any host/*ingress* conflict anyway when we come to check the other port.

XXX: host/host collisions need to be avoided by the scheduler, not the allocator. However:

"the scheduler is not involved in host mode ports. it was a very rushed feature, if *i* recall correctly, and it's sensitive to collisions."

$\land\ spec.mode = ingress \Rightarrow$
  $\forall\ eid2 \in \text{DOMAIN}\ endpoints :$
  $\forall\ i2 \in \text{DOMAIN}\ endpoints[eid2].ports :$
    $\langle eid,\ i\rangle \neq \langle eid2,\ i2\rangle \Rightarrow$ Don't check a port against itself
    LET $p2 \triangleq endpoints[eid2].ports[i2]$
    IN
      The other port must have a different network address:
      $\lor\ Addr(port) \neq Addr(p2)$
    XXX: an exception to this rule for *ingress*/host conflicts: We can't use the same address for an *ingress* and a host port because an *ingress* port must be allocated on every node, and so would conflict with the host port. However, this is a known bug in *SwarmKit*. For now, ignore *ingress*/host conflicts:
      $\lor\ p2.mode = host$

Check that *spec* is *OK* in itself (ignoring any other *endpoints*).
$SpecOK(spec) \triangleq$
  $\forall\ i \in \text{DOMAIN}\ spec :$ For every pair of ports $\langle i,\ j\rangle$
  $\forall\ j \in 1\,..\,(i-1) :$
    $\lor\ spec[i].mode = host$ Don't care about host-mode ports
    $\lor\ spec[j].mode = host$
    $\lor\ Addr(spec[i]) \neq Addr(spec[j])$ The requested addresses are different
    $\lor\ spec[i].published\_port = dynamic$ or they are both dynamic.

Special value to indicate creation of a new endpoint.
$nullId \triangleq \text{CHOOSE}\ x : x \notin EndpointID$

Checks that the allocator rejects a request only if it should.
$RejectJustified \triangleq$
  $\forall\ s \in EndpointSpec :$ s is the new spec
    $\forall\ eid \in \text{DOMAIN}\ endpoints \cup \{nullId\} :$ *eid* is the endpoint to update, or *nullId* for creation
    LET $oldEndpoint \triangleq \text{IF}\ eid = nullId\ \text{THEN}\ nullEndpoint\ \text{ELSE}\ endpoints[eid]$
      The possible allocations, or {} if rejected:
      $props \triangleq Allocate(oldEndpoint,\ s)$
      Ports used in our old configuration. We can conflict with these:
      $dealloc \triangleq \{Addr(p) : p \in \{p \in Range(oldEndpoint.ports) : p.mode = ingress\}\}$
      Ports not used by us:
      $usedByOthers \triangleq allocated \setminus dealloc$
      Ingress ports where the user chose the port number:
      $staticPorts \triangleq \{p \in Range(s) : \land\ p.mode = ingress$
                                                $\land\ p.published\_port \neq dynamic\}$
      All the *ingress* addresses chosen by the user:
      $staticAddr \triangleq \{Addr(p) : p \in staticPorts\}$
      We expect the allocation to be rejected:
      $rejectOK \triangleq$
        The specification is itself invalid:
        $\lor\ \neg SpecOK(s)$

```
                We asked for a port that is already in use:
            ∨ staticAddr ∩ usedByOthers ≠ {}
                There aren't enough free dynamic addresses for some protocol:
            ∨ ∃ proto ∈ Protocol :
                LET dynNeeded ≜ Cardinality({i ∈ DOMAIN s :
                                             ∧ s[i].protocol = proto
                                             ∧ s[i].mode = ingress
                                             ∧ s[i].published_port = dynamic})
                    dynAvail  ≜ Cardinality({a ∈ Address \ (usedByOthers ∪ staticAddr) :
                                             ∧ a[1] = proto
                                             ∧ a[2] ∈ DynamicPortNumber})
                Note: dynNeeded is an over-estimate because we might be able to reuse an existing
                static address.
                IN  dynNeeded > dynAvail
        IN
            If the allocator rejected the new spec, we understand why:
            ∨ props = {} ⇒ rejectOK
            ∨ Print(⟨props, rejectOK, endpoints, eid, s⟩, FALSE)    Log the reason on error

    If an endpoint's spec didn't change, then its allocation shouldn't change either. This tests that
    Allocate is idempotent. XXX: This is not currently the case, because if we have two similar
    specs then we only copy the existing allocation for the first one.
StepAllocateIdempotent ≜
    ∀ eid ∈ DOMAIN endpoints ∩ DOMAIN endpoints′ :
    LET ep ≜ endpoints[eid]
    IN
    ep.spec = (ep.spec)′
     ⇒
    ep.ports = (ep.ports)′

    All steps are idempotent.
AllocateIdempotent ≜
    □[StepAllocateIdempotent]_vars
```